

# Sequence Labeling

He He

New York University

October 6, 2021

## Last week

Goal: probabilistic modeling of (text) sequences

N-gram models

- ▶ Markov assumption: the next word depends on limited prior context
- ▶ Tackling sparsity
  - ▶ Discounting: allocate some probability mass to unseen events
  - ▶ backoff/interpolation: use dynamic context

LM as sequence classification

- ▶ Log-linear LM: represent context by a (handcrafted) feature vector
- ▶ Feed-forward neural LM: represent context by concatenated word vectors
- ▶ Recurrent neural LM: represent context by a recurrently updated state

# Table of Contents

Introduction

Maximum-entropy Markov Models

Conditional Random Field

Neural Sequence Modeling

# Sequence labeling

Language modeling as sequence labeling:


the	fox	jumped	over	the	dog	STOP
↑	↑	↑	↑	↑	↑	↑
*	the	fox	jumped	over	the	dog

**Part-of-speech (POS) tagging:**

DT	NN	VBD	IN	DT	NN
↑	↑	↑	↑	↑	↑
the	fox	jumped	over	the	dog

## Span prediction

### Named-entity recognition (NER):



### BIO notation:

- ▶ Reduce span prediction to sequence labeling
- ▶ B-<tag>: the first word in span <tag>
- ▶ I-<tag>: other words in span <tag>
- ▶ O: words not in any span

# POS tagging

**Part-of-speech:** the *syntactic* role of each word in a sentence

POS tagset:

- ▶ Universal dependency tagset
  - ▶ **Open class tags:** content words such as nouns, verbs, adjectives, adverbs etc.
  - ▶ **Closed class tags:** function words such as pronouns, determiners, auxiliary verbs etc.
- ▶ Penn Treebank tagset (developed for English, 45 tags)

Application:

- ▶ Often the first step in the NLP pipeline.
- ▶ Used as features for other NLP tasks.
- ▶ Included in tools such as Stanford CoreNLP and spaCy.

## The majority baseline

A dumb approach: look up each word in the dictionary and return the most common POS tag.

## The majority baseline

A dumb approach: look up each word in the dictionary and return the most common POS tag.

Problem: **ambiguity**. Example?

Types:		WSJ	Brown
Unambiguous	(1 tag)	44,432 (86%)	45,799 (85%)
Ambiguous	(2+ tags)	7,025 (14%)	8,050 (15%)
Tokens:			
Unambiguous	(1 tag)	577,421 (45%)	384,349 (33%)
Ambiguous	(2+ tags)	711,780 (55%)	786,646 (67%)

**Figure 8.2** Tag ambiguity for word types in Brown and WSJ, using Treebank-3 (45-tag) tagging. Punctuation were treated as words, and words were kept in their original case.

Most types are unambiguous, but ambiguous ones are common words!

Most common tag: 92% accuracy on WSJ (vs 97% SOTA)

Always compare to the majority class baseline.



# Table of Contents

Introduction

Maximum-entropy Markov Models

Conditional Random Field

Neural Sequence Modeling

## Multiclass classification

**Task:** given  $x = (x_1, \dots, x_m) \in \mathcal{X}^m$ , predict  $y = (y_1, \dots, y_m) \in \mathcal{Y}^m$ .

**Predictor:** *independent* classification problem at each position  $y_i = h(x, i) \quad \forall i$

Multinomial logistic regression ( $\theta \in \mathbb{R}^d$ ):

$$p(y_i | x) =$$

Feature templates:

$$T(x, i, y) =$$

## Multiclass classification

**Task:** given  $x = (x_1, \dots, x_m) \in \mathcal{X}^m$ , predict  $y = (y_1, \dots, y_m) \in \mathcal{Y}^m$ .

**Predictor:** *independent* classification problem at each position  $y_i = h(x, i) \quad \forall i$

Multinomial logistic regression ( $\theta \in \mathbb{R}^d$ ):

$$p(y_i | x) = \frac{\exp[\theta \cdot \phi(x, i, y_i)]}{\sum_{y' \in \mathcal{Y}} \exp[\theta \cdot \phi(x, i, y')]}$$

- ▶ Learning: MLE (is the objective convex?)
- ▶ Inference: trivial ( $\arg \max_{y \in \mathcal{Y}} p(y | x)$ )
- ▶ Does not consider dependency among  $y_i$ 's.

DT NN ?

B-<org> I-<org> ?

# Maximum-entropy markov model (MEMM)

Model the joint probability of  $y_1, \dots, y_m$ :

$$p(y_1, \dots, y_m \mid x) = \prod_{i=1}^m p(y_i \mid y_{i-1}, x) .$$

- ▶ Use the Markov assumption similar to n-gram LM.
- ▶ Insert start/end symbols:  $y_0 = *$  and  $y_m = \text{STOP}$ .

Parametrization:

$$p(y_i \mid y_{i-1}, x) = \frac{\exp [\theta \cdot \phi(x, i, y_i, y_{i-1})]}{\sum_{y' \in \mathcal{Y}} \exp [\theta \cdot \phi(x, i, y', y_{i-1})]}$$

Learning: MLE (each sequence produces  $m$  classification examples)

## Features for POS tagging

Interaction between word and tags:

$$\begin{aligned} &\mathbb{1}\{x_i = \textit{the}, y_i = \text{DET}\} \\ &\mathbb{1}\{y_i = \text{PROPN}, x_{i+1} = \textit{Street}, y_{i-1} = \text{NUM}\} \\ &\mathbb{1}\{y_i = \text{VERB}, y_{i-1} = \text{AUX}\} \end{aligned}$$

Word shape feature that help with unknown words:

$x_i$  contains a particular prefix (perhaps from all prefixes of length  $\leq 2$ )  
 $x_i$  contains a particular suffix (perhaps from all suffixes of length  $\leq 2$ )  
 $x_i$ 's word shape  
 $x_i$ 's short word shape

# Inference

## Decoding / Inference:

$$\begin{aligned} & \arg \max_{y \in \mathcal{Y}^m} \prod_{i=1}^m p(y_i \mid y_{i-1}, x) \\ &= \arg \max_{y \in \mathcal{Y}^m} \sum_{i=1}^m \log p(y_i \mid y_{i-1}, x) \\ &= \arg \max_{y \in \mathcal{Y}^m} \sum_{i=1}^m \underbrace{s(y_i, y_{i-1})}_{\text{local score}}, \end{aligned}$$

where  $s(y_i, y_{i-1}) = \theta \cdot \phi(x, i, y_i, y_{i-1})$ .

- ▶ Bruteforce: exact, compute scores of all sequences,  $O(|\mathcal{Y}|^m)$
- ▶ Greedy: inexact, predict  $y_i$  sequentially,  $O(m)$


## Viterbi decoding

$$\begin{aligned}& \max_{y \in \mathcal{Y}^m} \sum_{i=1}^m s(y_i, y_{i-1}) \\&= \max_{y \in \mathcal{Y}^m} \left( \sum_{i=1}^{m-1} s(y_i, y_{i-1}) + s(y_m, y_{m-1}) \right) \\&= \max_{y_m \in \mathcal{Y}} \max_{y \in \mathcal{Y}^{m-1}} \left( \sum_{i=1}^{m-1} s(y_i, y_{i-1}) + s(y_m, y_{m-1}) \right) \\&= \max_{y_m \in \mathcal{Y}} \max_{t \in \mathcal{Y}} \max_{y \in \mathcal{Y}^{m-1}, y_{m-1}=t} \left( \sum_{i=1}^{m-1} s(y_i, y_{i-1}) + s(y_m, y_{m-1}=t) \right) \\&= \max_{y_m \in \mathcal{Y}} \max_{t \in \mathcal{Y}} \left( s(y_m, t) + \max_{y \in \mathcal{Y}^{m-1}, y_{m-1}=t} \sum_{i=1}^{m-1} s(y_i, y_{i-1}) \right) \\&= \max_{y_m \in \mathcal{Y}} \max_{t \in \mathcal{Y}} \underbrace{(s(y_m, t) + \pi[m-1, t])}_{\pi[m, y_m]}\end{aligned}$$

## Viterbi decoding

DP:  $\pi[j, t] = \max_{t' \in \mathcal{Y}} \pi[j-1, t'] + s(y_j = t, y_{j-1} = t')$

Backtracking:  $p[j, t] = \arg \max_{t' \in \mathcal{Y}} \pi[j-1, t'] + s(y_j = t, y_{j-1} = t')$



Time complexity?



# Summary

Sequence labeling:  $\mathcal{X}^m \rightarrow \mathcal{Y}^m$

- ▶ **Majority baseline:**  $y_i = h(x_i)$  (no context)
- ▶ **Multiclass classification:**  $y_i = h(x, i)$  (global input context)
- ▶ **MEMM:**  $y_i = h(x, i, y_{i-1})$  (global input context, previous output context)

Problem:  $y_t$  cannot be influenced by future evidence (more on this later)

Next: score  $x$  and the output  $y$  instead of local components  $y_i$

# Table of Contents

Introduction

Maximum-entropy Markov Models

Conditional Random Field

Neural Sequence Modeling

# Structured prediction

**Task:** given  $x = (x_1, \dots, x_m) \in \mathcal{X}^m$ , predict  $y = (y_1, \dots, y_m) \in \mathcal{Y}^m$ .

- ▶ Similar to multiclass classification except that  $\mathcal{Y}$  is very large
- ▶ Compatibility score:  $h: \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$
- ▶ Predictor:  $\arg \max_{y \in \mathcal{Y}^m} h(x, y)$

General idea:

- ▶  $h(x, y) = f(\theta \cdot \Phi(x, y))$
- ▶  $\Phi$  should be decomposable so that inference is tractable
- ▶ Loss functions: structured hinge loss, negative log-likelihood etc.
- ▶ Inference: viterbi, interger linear programming (ILP)

# Graphical models

Graphical model:

- ▶ A joint distribution of a set of random variables
- ▶ A graph represents conditional independence structure among random variables
  - ▶ Nodes: random variables
  - ▶ Edges: dependency relations
- ▶ Learning: estimate parameters of the distribution from data
- ▶ Inference: compute conditional/marginal distributions

# Directed graphical model

Directed graphical model (aka Bayes nets):

- ▶ Edges represent conditional dependencies

Example: MEMM

# Undirected graphical models

Undirected graphical model (aka Markov random field):

- ▶ More natural for relational or spatial data

**Conditional random field:**

- ▶ MRF conditioned on observed data
- ▶ Parameterization:

$$p(y \mid x; \theta) = \frac{1}{Z(x, \theta)} \prod_{c \in \mathcal{C}} \psi_c(y_c \mid x; \theta)$$

- ▶ Clique  $\mathcal{C}$ : a subset of nodes/variables that form a complete graph
- ▶  $\psi_c$ : non-negative clique potential functions, also called factors
- ▶  $Z(x, \theta)$ : partition function (normalizer)

## Linear-chain CRF

Model dependence among  $Y_i$ 's

$$p(y \mid x; \theta) = \frac{1}{Z(x, \theta)} \prod_{i=1}^m \psi_i(y_1, \dots, y_m \mid x; \theta)$$

## Linear-chain CRF

Model dependence among *neighboring*  $Y_i$ 's

$$p(y \mid x; \theta) = \frac{1}{Z(x, \theta)} \prod_{i=1}^m \psi_i(y_i, y_{i-1} \mid x; \theta)$$



## Linear-chain CRF for sequence labeling

Log-linear potential function:

$$\psi_i(y_i, y_{i-1} \mid x; \theta) = \exp(\theta \cdot \phi(x, i, y_i, y_{i-1}))$$

$$\begin{aligned} p(y \mid x; \theta) &\propto \prod_{i=1}^m \exp(\theta \cdot \phi(x, i, y_i, y_{i-1})) \\ &= \exp\left(\sum_{i=1}^m \theta \cdot \phi(x, i, y_i, y_{i-1})\right) \end{aligned}$$

Log-linear model with decomposable global feature function:

$$\begin{aligned} \Phi(x, y) &\stackrel{\text{def}}{=} \sum_{i=1}^m \phi(x, i, y_i, y_{i-1}) \\ p(y \mid x; \theta) &= \frac{\exp(\sum_{i=1}^m \theta \cdot \phi(x, i, y_i, y_{i-1}))}{\sum_{y' \in \mathcal{Y}^m} \exp(\sum_{i=1}^m \theta \cdot \phi(x, i, y'_i, y'_{i-1}))} \\ &= \frac{\exp(\theta \cdot \Phi(x, y))}{\sum_{y' \in \mathcal{Y}^m} \exp(\theta \cdot \Phi(x, y'))} \end{aligned}$$

# Learning

MLE:

$$\begin{aligned}\ell(\theta) &= \sum_{(x,y) \in \mathcal{D}} \log p(y \mid x; \theta) \\ &= \sum_{(x,y) \in \mathcal{D}} \log \frac{\exp(\theta \cdot \Phi(x, y))}{\sum_{y' \in \mathcal{Y}^m} \exp(\theta \cdot \Phi(x, y'))}\end{aligned}$$

- ▶ Is the objective differentiable?
- ▶ Use back-propagation (autodiff) (equivalent to the forward-backward algorithm).
- ▶ Main challenge: compute the partition function.

$$\begin{aligned}
& \log \sum_{y \in \mathcal{Y}^m} \exp \left( \sum_{i=1}^m s(y_i, y_{i-1}) \right) \\
&= \log \sum_{y \in \mathcal{Y}^m} \left( \exp \left( \sum_{i=1}^{m-1} s(y_i, y_{i-1}) + s(y_m, y_{m-1}) \right) \right) \\
&= \log \sum_{y_m \in \mathcal{Y}} \sum_{t \in \mathcal{Y}} \sum_{y \in \mathcal{Y}^{m-1}, y_{m-1}=t} \exp \left( \sum_{i=1}^{m-1} s(y_i, y_{i-1}) + s(y_m, y_{m-1}=t) \right) \\
&= \log \sum_{y_m \in \mathcal{Y}} \sum_{t \in \mathcal{Y}} \exp(s(y_m, y_{m-1}=t)) \sum_{y \in \mathcal{Y}^{m-1}, y_{m-1}=t} \exp \left( \sum_{i=1}^{m-1} s(y_i, y_{i-1}) \right) \\
&= \log \sum_{y_m \in \mathcal{Y}} \sum_{t \in \mathcal{Y}} \exp(s(y_m, y_{m-1}=t)) \exp(\pi[m-1, t]) \\
&= \log \sum_{y_m \in \mathcal{Y}} \sum_{t \in \mathcal{Y}} \underbrace{\exp(s(y_m, y_{m-1}=t) + \pi[m-1, t])}_{\exp(\pi[m, y_m])}
\end{aligned}$$

## Compute the partition function

$$\exp(\pi[j, t]) \stackrel{\text{def}}{=} \sum_{y \in \mathcal{Y}^j, y_j = t} \exp \left( \sum_{i=1}^j s(y_i, y_{i-1}) \right)$$

$$\pi[j, t] \stackrel{\text{def}}{=} \log \sum_{y \in \mathcal{Y}^j, y_j = t} \exp \left( \sum_{i=1}^j s(y_i, y_{i-1}) \right)$$

$$\pi[j, t] = \log \sum_{t' \in \mathcal{Y}} \exp (s(y_j = t, y_{j-1} = t') + \pi[j-1, t'])$$

Compute the partition function

## Compute the partition function

DP:

$$\exp(\pi[j, t]) = \sum_{t' \in \mathcal{Y}} \exp(s(y_j = t, y_{j-1} = t') + \pi[j-1, t'])$$

$$\pi[j, t] = \log \sum_{t' \in \mathcal{Y}} \exp(s(y_j = t, y_{j-1} = t') + \pi[j-1, t'])$$

*The logsumexp function:*

$$\text{logsumexp}(x_1, \dots, x_n) = \log(e^{x_1} + \dots + e^{x_n})$$

$$\text{logsumexp}(x_1, \dots, x_n) = x^* + \log(e^{x_1 - x^*} + \dots + e^{x_n - x^*})$$

- ▶ Same as Viterbi except that max is replaced by logsumexp.
- ▶ Is this a coincidence?

$$\max(a + b, a + c) = a + \max(b, c)$$

$$\text{logsumexp}(a + b, a + c) = a + \text{logsumexp}(b, c)$$

## Learning

Use forward algorithm to compute:

$$\text{loss} = -\ell(\theta, x, y) = -\log \frac{\exp(\theta \cdot \Phi(x, y))}{\sum_{y' \in \mathcal{Y}^m} \exp(\theta \cdot \Phi(x, y))}$$

`loss.backward()`

Exercise: show that the optimal solution satisfies

$$\sum_{(x,y) \in \mathcal{D}} \Phi_k(x, y) = \sum_{(x,y) \in \mathcal{D}} \mathbb{E}_{y \sim p_\theta} [\Phi_k(x, y)]$$

Interpretation: Observed counts of feature  $k$  equals expected counts of feature  $k$ .

# Inference

$$\begin{aligned} & \arg \max_{y \in \mathcal{Y}^m} \log p(y \mid x; \theta) \\ &= \arg \max_{y \in \mathcal{Y}^m} \log \exp(\theta \cdot \Phi(x, y)) - \log Z(\theta) \\ &= \arg \max_{y \in \mathcal{Y}^m} \sum_{i=1}^m s(y_i, y_{i-1}) \end{aligned}$$

- ▶ Find highest-scoring sequence.
- ▶ Use Viterbi + backtracking.



# Summary

## Conditional random field

- ▶ Undirected graphical model
- ▶ Use factors to capture dependence among random variables
- ▶ Need to trade-off modeling and inference

## Linear-chain CRF for sequence labeling

- ▶ Models dependence between neighboring outputs
- ▶ Learning: forward algorithm + backpropagation
- ▶ Inference: Viterbi algorithm

# Table of Contents

Introduction

Maximum-entropy Markov Models


Conditional Random Field

Neural Sequence Modeling

# Classification using recurrent neural networks

Logistic regression with  $h_t$  as the features:

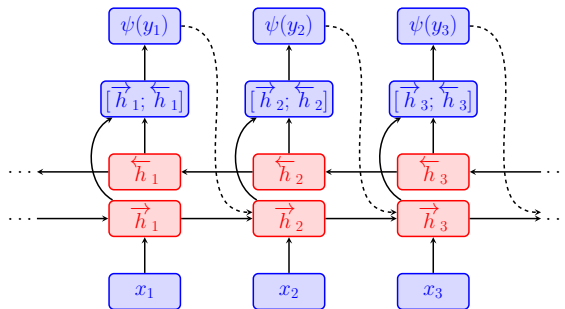
$$p(y_i | x) = \text{softmax}(W_{ho}h_i + b)$$



What is the problem?

# Bi-directional RNN

Use two RNNs to summarize the “past” and the “future”:



- ▶ Concatenated hidden states:  $h_i = [\vec{h}_{1:m}; \overleftarrow{h}_{1:m}]$
- ▶ Optional: use  $y_{i-1}$  as inputs:  $\vec{h}'_i = [\vec{h}_i; \underbrace{W_{yh}y_{i-1}}_{\text{label embedding}}]$

## Bi-LSTM CRF

Use neural nets to compute the local scores:

$$s(y_i, y_{i-1}) = s_{\text{unigram}}(y_i) + s_{\text{bigram}}(y_i, y_{i-1})$$

Basic implementation:

$$\begin{aligned} s_{\text{unigram}}(y_i) &= (W_{ho} h_i + b)[y_i] \\ s_{\text{bigram}}(y_i, y_{i-1}) &= \theta_{y_i, y_{i-1}} \quad (|\mathcal{Y}|^2 \text{ parameters}) \end{aligned}$$

Context-dependent scores:

$$\begin{aligned} s_{\text{unigram}}(y_i) &= (W_{ho} h_i + b)[y_i] \\ s_{\text{bigram}}(y_i, y_{i-1}) &= w_{y_i, y_{i-1}} \cdot h_i + b_{y_i, y_{i-1}} \end{aligned}$$

# Does it worth it?

Typical neural sequence models:

$$p(y \mid x; \theta) = \prod_{i=1}^m p(y_i \mid x, y_{1:i-1}; \theta)$$

*Exposure bias*: a learning problem

- ▶ Conditions on gold  $y_{1:i-1}$  during training but predicted  $\hat{y}_{1:i-1}$  during test
- ▶ Solution: search-aware training

*Label bias*: a model problem

- ▶ Locally normalized models are strictly less expressive than globally normalized *given partial inputs* [Andor+ 16]
- ▶ Solution: globally normalized models or better encoder

## Does it worth it?

Empirical results from [Goyal+ 19]

	Unidirectional	Bidirectional
pretrain-greedy	76.54	92.59
pretrain-beam	77.76	93.29
locally normalized	83.9	<b>93.76</b>
globally normalized	<b>83.93</b>	93.73

Table 2: **Accuracy results on CCG supertagging when initialized with a regular teacher-forcing model.** Reported using *Unidirectional* and *Bidirectional* encoders respectively with fixed attention tagging decoder. *pretrain-greedy* and *pretrain-beam* refer to the output of decoding the initializer model. *locally normalized* and *globally normalized* refer to search-aware soft-beam models